

DAIIA Assignment Report

Course: [Distributed Artificial Intelligence and AI Agents]

Assignment: Homework 3

Lorenzo Deflorian, Riccardo Fragale, Juozas Skarbalius

KTH Royal Institute of Technology

November 28, 2025

Contents

Running Instructions

Import the ZIP file given on GAMA and go to the folder assignment3/models. You will find a file called NQueen.gaml where our model for task 1 is implemented. Then, from the interface of GAMA, click on the play button and you will see a simulation. Use the tools provided by the GAMA simulation interface to adjust the speed, read the outputs and verify on screen that everything is working correctly. There is a parameter in the global of the module, called queens, that can be modified and defines the number of queens in the chess board. It is modifiable and should be used to verify that our code works for all N between 4 and 19. Regarding task 1, in the same folder assignment3/models there is a .gaml file called FestivalStages. Also in this case open it and play the simulation.

1. Task1: NQueen

General overview

We have to solve the N Queens problem; in particular the following are the rules of the game:

- Create a $N \times N$ size chessboard, placing N queens on it
- No two queens can share the same row
- No two queens can share the same column
- No two queens can share the same diagonal line

Starting from a random situation the Queens should adjust their position so that they do not violate the rules and they found a correct arrangement. This property must be valid for $N \in \{4, \dots, 19\}$. In our setup we have a ChessBoard that appears on screen and a certain number of Queen agents, defined by the variable `queens` (as said in the running instructions). Each queen is an agent, they are able to communicate between each other through "*fipa-contract-net*" protocol. In our solution each queen is only able to communicate to its predecessor and to its successor. There is a sort of recursive procedure in which if a queen has no available position, she must let her predecessor know and ask her to reposition her. If also the predecessor has no available positions left, she must message her predecessor and so on and so for up until a correct arrangement is found. Queens in correct positions are depicted in green while queens that have not found the correct placement are depicted in red.

Code

First of all, each Queen is positioned on the chessBoard almost randomly using the following init procedure:

```
chessBoardCell myCell <- one_of (chessBoardCell);
```

```
[...]
init {
    //Assign a free cell
    loop cell over: myCell.neighbours{
        if cell.queen = nil{
            myCell <- cell;
            break;
        }
    }

    location <- myCell.location;
    myCell.queen <- self;
    add self to: allQueens;
    do refreshOccupancyGrid;
}
```

The procedure starts by calculating the occupancyGrid and verifying whether there are conflicts with respect to the rules. When a conflict is found, queens need to move. The core relocation logic lives in the needToMove action. If a queen is in a conflict (indicated by a red color) and has no immediately free safe cells, it initiates negotiation. Unlike the previous approach based on visibility, the queens are organized in a circular linked list. Each queen is assigned a specific `predecessor` and `successor` during initialization. When stuck, a queen sends a request specifically to its predecessor.

```
// Communicate via chain: ONLY ask predecessor
if predecessor != nil and !awaitingResponse{
    do start_conversation to: [predecessor]
    protocol: 'fipa-contract-net' performative: 'cfp'
    contents: ['request_move', string(self.myCell.grid_x),
               string(self.myCell.grid_y), name];

    awaitingResponse <- true;
    messageContext <- "chain_request";
}
```

The requester includes its own name in the message so that if the request is forwarded down the chain, the eventual helper knows who to reply to.

Incoming CFPs trigger the `handleCFP` reflex. The logic follows a "Chain of Responsibility" pattern:

1. **If the receiver can move:** It moves to a free spot, effectively freeing its old cell. It then sends a PROPOSE message directly to the `originalRequester` (extracted from the message content) telling them the position is now available.
2. **If the receiver cannot move:** It acts as a middleman. It forwards the CFP to its own `predecessor`, preserving the `originalRequester`'s name. It then sends a REFUSE to the immediate sender to close that specific transaction.

- 3. Chain Exhaustion:** If the request circles back to the original sender (Full Circle), the chain is exhausted, and the request is dropped to prevent infinite loops.

```

reflex handleCFP when: !empty(cfps){
    if length(myOptions) > 0 {
        // Move and notify original requester
        do start_conversation to: [originalQueen] ...
        performative: 'propose'
        contents: ['position_available', string(oldCell.grid_x), ...];
    } else {
        // Forward to predecessor
        do start_conversation to: [predecessor] ... performative: 'cfp'
        contents: ['request_move', ..., originalRequester];
    }
}

```

The `handlePropose` reflex processes the resolution. When a queen in the chain finds a spot and moves, it sends a '`position_available`' message to the queen that started the chain.

```

reflex handlePropose when: !empty(proposes) and awaitingResponse{
    if contents[0] = 'position_available'{
        // Move to the cell freed by the helper
        myCell <- freedCell;
        location <- freedCell.location;
        do accept_proposal
        message: proposalMessage contents: ['move_completed'];
    }
}

```

This creates a robust system where a single request can ripple through the entire population of queens until one is found that can shift to accommodate the group, resolving deadlocks that occurred in the visibility-based approach.

Finally, to provide visual feedback on the solution state, the queens update their color dynamically:

- **Red:** The queen is currently in a conflict (occupancy grid value > 0).
- **Green:** The queen is in a safe position (occupancy grid value $== 0$).

The simulation is considered solved when all queens turn green. The result is clearer separation: messages carry intent, reflexes react to performatives, and actions perform local state updates.

Demonstration

We will show below a couple of situations proving that our implementation is correct using $N = 4$. In our opinion it would definitely be better for a reader to run many simulations and verify that the N Queen problem is solved for that randomical beginning position.

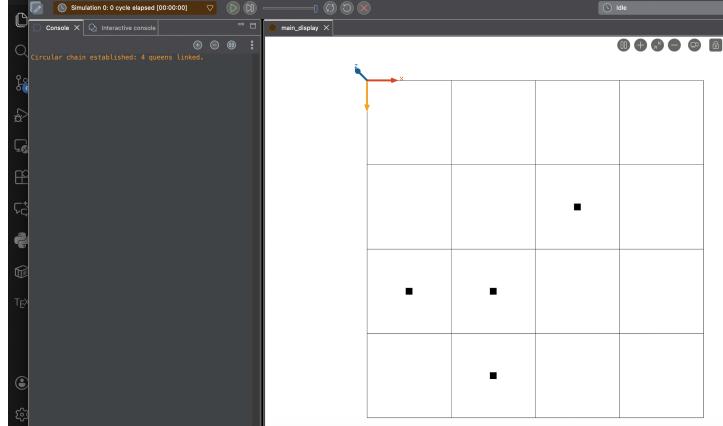


Figure 1: Starting point

As it can be seen from ?? the queens are placed correctly on the map.

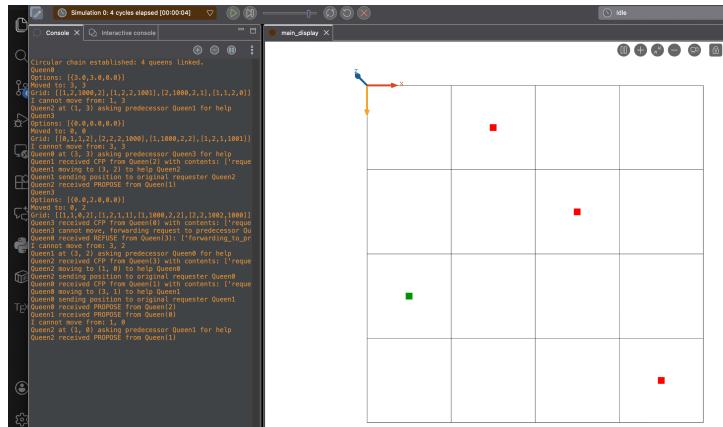


Figure 2: Queens are moving

?? shows that queens are communicating and changing their position in order to find a correct placement. It is clear also that sometimes the request is passed from predecessor to predecessor until someone can move.

Finally, ?? shows that the position is found.

We will also add a picture of what happens inside a run with $N=19$. Clearly the logs are more and the situation might be definitely more complex so more loops are needed to find the correct placement for all the queens in the chessboard.

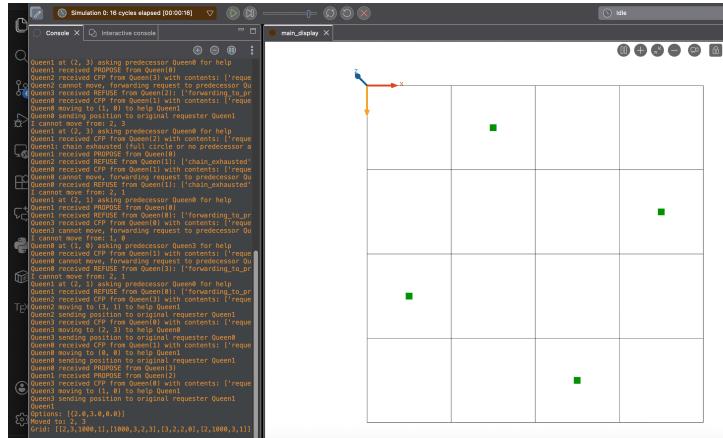


Figure 3: Position found

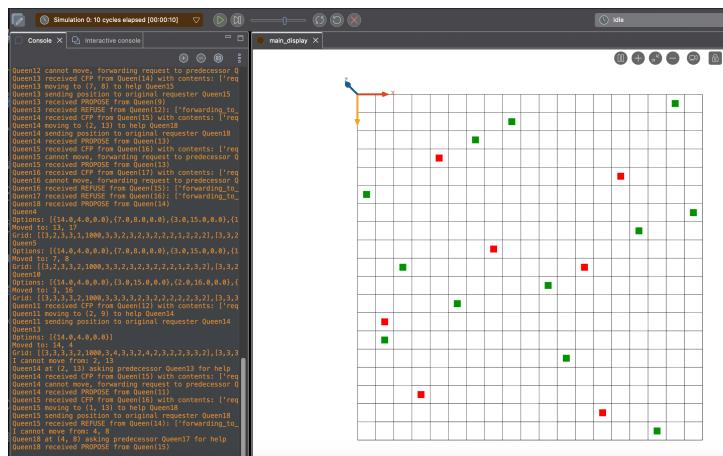


Figure 4: 19 queens moving on the chessboard

2. Task 2: Positioning speakers at main stage

For the purpose of this task we decide to eliminate the features of the previous homeworks related to the festival except for the InformationCenter and the skeleton of the normal guest. We were asked to put some stages into the festival which host an act for a different amount of time. The guest must know their position since the beginning of the simulation and must decide to go and attend an event at a festival based on a utility function. Basically, the decision is based on the guest preference and the characteristic of each stage related to some properties (e.g sound system ,light visuals etc). Guest communicate with stages through FIPA and after a "discussion" they must select the stage with the highest utility.

3.1 Explanation

Describe what this section requires and how you approached solving it.

3.2 Code

During initialization the model creates one InformationCenter, three Stage agents and ten Guests. Each guest receives a list of all stages (and also knows their position) as shown below

```
ask guests {
    set stages <- allStages;
}
```

Each Stage has three key properties:

- lightShow (0–100)
- speaker (0–100)
- musicStyle (0–100)

These are randomly generated and represent the stage's entertainment characteristics. Stages also have a starting time and a duration, both different for each stage and randomly generated. Each guest has individual preferences (lightShowPreference, speakerPreference, musicStylePreference), all randomly initialized in the interval 0 to 100. Utility for a stage is computed as follows:

```
action compute_utilities(list<int> stats, Stage sender) {
    int lightShow <- stats[0];
    int speaker <- stats[1];
    int musicStyle <- stats[2];
    int utility <- lightShowPreference * lightShow
    + speakerPreference * speaker + musicStylePreference * musicStyle;
    utilities[sender] <- utility;
}
```

Guests request stage statistics using FIPA Contract Net Protocol. The guest sends a cfp (Call For Proposal) performative to all stages with contents ["stats"].

```
reflex request_utilities when: hasRequestedUtilities = false {
    write 'Requesting utilities from stages';
    do start_conversation to: list(Stage)
    protocol: 'fipa-contract-net' performative: 'cfp' contents: ["stats"];
    hasRequestedUtilities <- true;
}
```

When a stage receives a CFP with content "stats", it responds immediately with a propose performative containing its three attributes as a list: [lightShow, speaker, musicStyle].

```

reflex receiveCFP when: !empty(cfps) {
    loop cfpMsg over: cfps {
        list contents_list <- list(cfpMsg.contents);
        if (contents_list[0] = "stats") {
            do propose message: cfpMsg contents:
                [lightShow, speaker, musicStyle];
        }
    }
}

```

All proposals are received in the receiveProposal reflex. For each one this reflex calls compute_utilities to score the stage and stores the result in a utilities map.

```

reflex receiveProposals when: hasReceivedUtilities = false and
hasRequestedUtilities = true and !empty(proposes){
    loop proposeMsg over: proposes {
        list<int> contents_list <- list<int>(list(proposeMsg.contents));
        do compute_utilities(contents_list, proposeMsg.sender);
    }
    hasReceivedUtilities <- true;
    do select_stage;
}

```

After processing all proposals, the guest invokes select_stage, which iterates through all stages and finds the one with the highest utility score stored in the utilities map.

```

action select_stage {
    int maxUtility <- 0;
    Stage bestStage <- nil;
    loop stage over: stages {
        if (utilities != nil and utilities contains_key(stage)) {
            int stageUtility <- utilities[stage];
            if (stageUtility > maxUtility) {
                maxUtility <- stageUtility;
                bestStage <- stage;
            }
        }
    }
    selectedStage <- bestStage;
}

```

Once a stage is selected and the current time falls within [startFestivalTime, endFestivalTime], the guest computes a random target location near the stage and moves toward it using the moving skill with movingSpeed = 0.75. Guests navigate to their selected stage, arrive at a random position nearby (within 6–12 units from the stage), and then dance continuously

3.3 Demonstration

I will show below screenshots for two cases;

- Guests request utilities to stages
 - Guest select the stage with higher utility

For the purpose of this demonstration I included only 4 guests so that the number of logging messages is reduced.

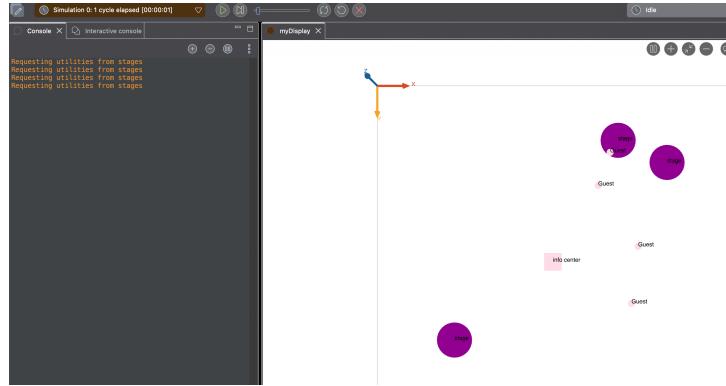


Figure 5: Guest request utilities

Guests request utilities to stages. A message is send to the stages in ?? to request utilities values.

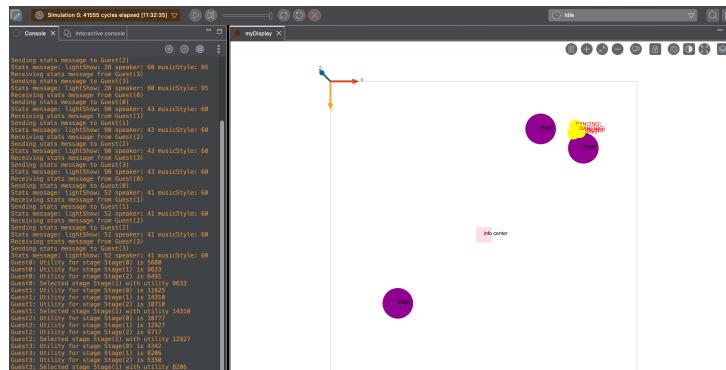


Figure 6: Guest select the stage with higher utility

Guest select the one with maximum utility. As it can be seen in ?? each guest is selecting the stage with the highest utility.

3. Challenge: Crowd Mass and Maximum Utility

The festival simulation described previously was further extended to include a crowd mass into a decision-making process for guests selecting stages to attend. Guests are as-

signed with random crowd mass preference (further called c_m), affecting their sensitivity to crowd sizes when evaluating stages. If such preference is high ($c_m > 50$), then guests prefer crowded stages, otherwise they prefer less crowded ones. The boundaries of c_m : $c_m \in [0, 100]$.

Furthermore, the utility function was modified to account for crowd mass at each stage, influencing guest initial and subsequent choices. First, the crowd mass is taken into account while each guest is making their initial choice of stage to attend. Then, guests exchange their initial stage choices with one another using FIPA protocol. Based on other guests decisions, each guest recalculates its crowd mass component of utility. Such changed crowd mass utility may lead to a different stage choice, which guests then proceed to attend. Moreover, Information Center agent acts as a leader and has all decision information from the guest agents. This allows Information Center to:

- Calculate the global utility of all guest agents.
- Issue proposals to specific agents to change their assigned stage in order to maximize overall utility.

4.1 Explanation. As mentioned, each guest is assigned a random crowd mass preference c_m at initialization. Let's begin by explaining what happens inside each individual agent. The initial behavior is not different from the previous task: each guest requests stage statistics using FIPA Contract Net Protocol and computes utilities for each stage. c_m is not taken into account at this particular point because it is most likely equal to zero as other agents are also performing this step at the same time. Instead, we assume that all stages might be occupied equally following normal distribution. This allows us to calculate the utility in a more fair way. The crowd mass utility then can be calculated as follows:

```

float assumedRatio <- (nbStages > 0) ? (1.0 / (nbStages as float)) : 0.0;
float preferredRatio <- (crowdMassPreference as float) / 100.0;

float diffAssumed <- assumedRatio - preferredRatio;
if (diffAssumed < 0.0) { diffAssumed <- -diffAssumed; }
float crowdU_assumed <- 1.0 - diffAssumed;
if (crowdU_assumed < 0.0) { crowdU_assumed <- 0.0; }

```

Afterwards, the local utility for each stage is calculated as:

$$(0.7 * (\text{baseU as float})) + (0.3 * \text{crowdU_assumed} * 1000000.0);$$

Where baseU is the utility of other stage characteristics (light show, speaker quality, music style), and the crowd mass utility is weighted at 30% of the total utility. Finally, the stage with the highest utility is selected as the initial choice, and leader with other guest agents are informed about this decision.

Subsequently, upon receiving other guest's choices, we can recalculate the c_m component and readjust the our stage choice if necessary. Additionally, in the meantime, the leader (Information Center) collects all guest decisions and computes the global utility of the system. If the leader identifies that the global utility can be improved by suggesting certain guests to change their stage, it sends proposals to those specific agents. Upon receiving such proposal, the guest simply updates its selected stage and proceeds to attend it. Such mechanism allows each guest agent to make their own decisions based on other agent's choices, while following orders from the leader in cases where global utility can be maximised.

To avoid infinite loops and freezes, the leader only suggests a single proposal (and together recalculates global utility) every 30th tick of the simulation, while guest will react to those proposals every 15th tick.

4.2 Code. Figure ?? illustrates the calculation of the global utility which is done by calculating each guests utility and summing them up. Such calculation might take place for up to 25 times during the simulation to ensure that the global utility is maximized as much as possible. Otherwise, the guest agents might maximize their own local utility without considering the overall utility of the environment.

```
// compute global utility
loop g over: allGuests {
    Stage st <- currentAssignment[g];
    if (st != nil and g.utilities != nil and g.utilities contains_key(st)) {
        int baseU <- g.utilities[st];

        int crowdOnStage <- currentCrowd[st];

        float crowdRatio <- 0.0;
        if (totalGuests > 0) {
            crowdRatio <- (crowdOnStage as float) / (totalGuests as float);
        }

        float preferredRatio <- (g.crowdMassPreference as float) / 100.0;

        // crowdU ∈ [0,1], maximum when crowdRatio == preferredRatio
        float diff <- crowdRatio - preferredRatio;
        if (diff < 0.0) {
            diff <- -diff;
        }

        float crowdU <- 1.0 - diff;
        if (crowdU < 0.0) {
            crowdU <- 0.0;
        }

        // weight crowd vs base: 70% base, 30% crowd
        float effective <- (0.7 * (baseU as float)) + (0.3 * crowdU * 1000000.0);
        newGlobalUtility <- newGlobalUtility + effective;
    }
}
currentGlobalUtility <- newGlobalUtility;
```

Figure 7: Global utility calculation by Information Center

Figure ?? illustrates the proposal mechanism used by the leader to suggest specific guest agents to change their selected stage. It works by evaluating the possible changes in global utility if a specific guest left or joined currently iterated stage. If such delta is higher than previously recorded maximum delta, the leader stores this information to be used later when sending proposals. Finally, after iterating through all stages and guests, the leader sends a proposal to the guest which would lead to the highest increase in global utility. Such mechanism is repeated every 30th tick of the simulation.

```

int totalGuests <= length(allGuests);
if (totalGuests == 0) {
    globalOptimumReached <- true;
} else {
    Guest bestGuest <- nil;
    Stage bestNewStage <- nil;
    float bestDeltaGlobal <- 0.0;
    // Try each guest as a candidate for switching
    loop g over: allGuests {
        Stage currentSt <- currentAssignment[g];
        // need a current assignment and a utilities map for this guest
        if (currentSt == nil and g.utilities containsKey(currentSt) and !g.optimizationFinished) {
            // ----- current effective utility of g on its current stage -----
            int baseCurrentG <- g.utilities[currentSt];
            int crowdCurrentA <- currentCrowd[currentSt];
            float crowdRatioCurrentA <- 0.0;
            if (totalGuests > 0) {
                crowdRatioCurrentA <- (crowdCurrentA as float) / (totalGuests as float);
            }
            float preferredRatioG <- (g.crowdMassPreference as float) / 100.0;
            float diffCurrentA <- crowdRatioCurrentA - preferredRatioG;
            if (diffCurrentA < 0.0) {
                diffCurrentA <- -diffCurrentA;
            }
            float crowdUG_A <- 1.0 - diffCurrentA;
            if (crowdUG_A < 0.0) {
                crowdUG_A <- 0.0;
            }
            float effG_A <- (0.7 * (baseCurrentG as float)) + (0.3 * crowdUG_A * 1000000.0);
            // ----- try moving g to each other stage B -----
            loop B over: allStages {
                ...
            }
        }
        // After checking all candidates, propose the single best switch (if any)
        if (bestGuest == nil and bestNewStage == nil and bestDeltaGlobal > 0.0) {
            write "Leader: proposing switch (GLOBAL) " + bestGuest + " from "
            + currentAssignment[bestGuest] + " to " + bestNewStage
            + " (" + (global! + bestDeltaGlobal + ")");
            do start_conversation
            to: [bestGuest]
            protocol: "fipa-contract-net"
            performative: "inform"
            contents: ["switch_to", bestNewStage];
        } else {
            // move increase global utility; count one "no improvement" step
            noImprovementSteps <- noImprovementSteps + 1;
            if (noImprovementSteps >= maxNoImprovementSteps and !globalOptimumReached) {
                globalOptimumReached <- true;
                write "***** MAX GLOBAL UTILITY APPROXIMATELY REACHED (no improving moves) *****";
                write "Best global utility: " + bestGlobalUtility;
                loop g over: allGuests {
                    do start_conversation
                    to: [g]
                    protocol: "fipa-contract-net"
                    performative: "inform"
                    contents: ["optimization_finished"];
                }
            }
        }
    }
}

```

Figure 8: Leader proposal mechanism (simplified view, details regarding delta utility calculation has been omitted for clarity)

4.3 Demonstration. As for the demonstration of this challenge, we will provide the following:

1. Figure ???: The initial environment setup where guests who prefer large crowds are colored in green, while those who prefer smaller crowds are colored in blue. The exact preference of crowd mass is also visible near each agent's position.
2. Figure ???: State after each guest has made their initial stage choices with text indicating that. Also initial utilities are calculated at this phase of execution.
3. Figure ???: State where global utility is reached and optimization is complete. The differences between initial and final stage choices can be visible in the figure provided below.
4. Figure ???: Final phase where each guest is enjoying the festival.

4. Final Remarks

In this homework we learnt a couple of things related to GAMA language and agents coordination. First of all, we were forced to create a "chain" of messages between agents that were bouncing back messages one to another (in the first task). Regarding the second task and its challenge it was not so easy to find a way to have a leader and to have the



Figure 9: Initial environment setup with crowd mass preferences

agents working together towards an objective following a utility function. We are really happy about our result and we are really looking forward to the final project which will certainly be more challenging.

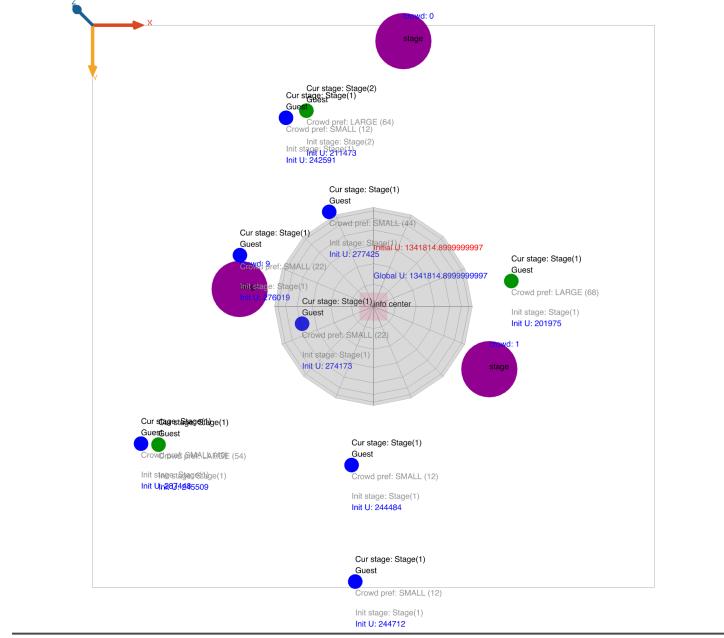


Figure 10: State after each guest has made their initial stage choices with text indicating that. Also initial utilities are calculated at this phase of execution.

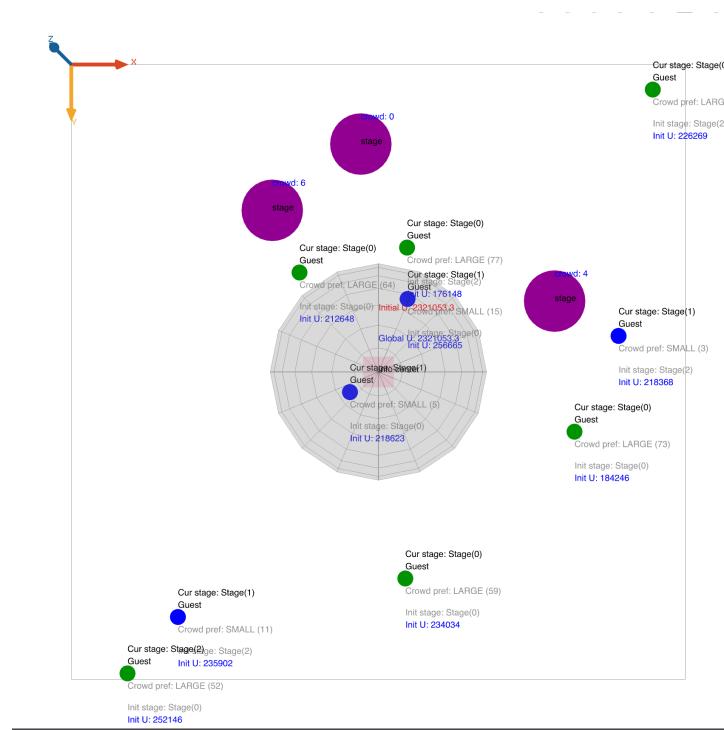


Figure 11: Leader imposes custom proposals to each agent to maximize global utility.

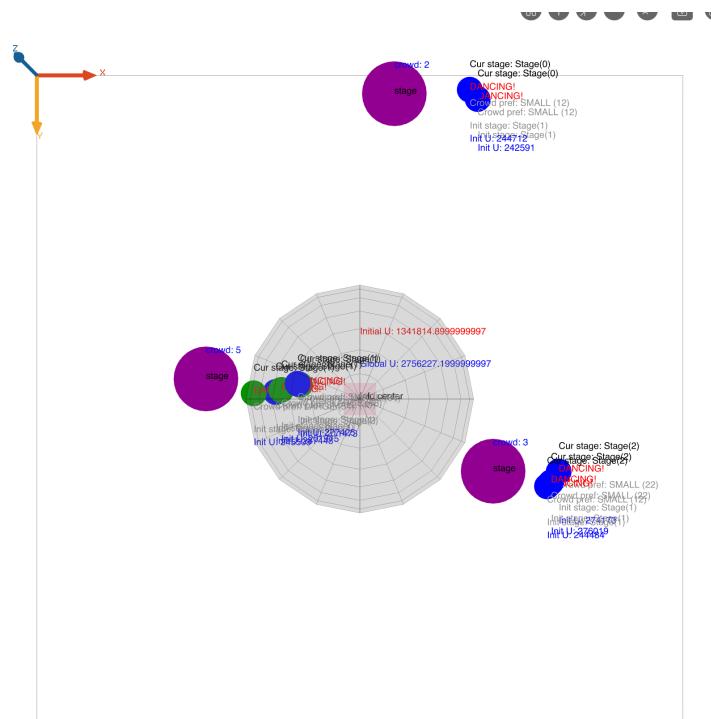


Figure 12: Each guest is enjoying the festival at their selected stage with the global utility maximized.